

Table E2: Support for Main-Paper Results across Alternative Mappings.

This table reports, for each model–prompt setup, the percentage of the N possible distinct maps that support each main-paper result. We count each unique map once. A map supports a result if its estimate matches the corresponding main-paper estimate in sign and in statistical significance (whether the estimate is statistically significant at the 10% level) or favors our hypothesis more strongly. GPT-based mapping uses GPT-6.1 Sol and Claude-based mapping uses Claude Sonnet 5.5. Prompt A supplies no investor background, and Prompt B adds the background of an investor who adopts the heuristic. As in the main paper, belief accuracy uses the two-sided wild-cluster bootstrap p -value by country, and delivery uses the one-sided p -value against equally sized random draws from the alternative heuristics.

Main-paper result	Analysis	Objective	GPT-based mapping		Claude-based mapping	
			Prompt A	Prompt B	Prompt A	Prompt B
			$N = 144$	$N = 360$	$N = 2,340$	$N = 5,616$
Figure 8	Belief accuracy	Overall	100%	100%	100%	100%
		Returns	100%	100%	100%	100%
		Upside	100%	100%	100%	100%
		Stability	100%	100%	100%	100%
Table 7	Historical delivery	Overall	100%	100%	100%	100%
		Returns	100%	100%	100%	100%
		Upside	100%	100%	100%	100%
		Stability	100%	100%	100%	100%
Table 8	Out-of-sample delivery	Overall	100%	100%	100%	100%
		Returns	100%	100%	87%	83%
		Upside	100%	100%	100%	100%
		Stability	89%	89%	71%	58%